

# DDA\_Tutorial\_1

April 29, 2022

Farjad Ahmed - 1747371 Tutorial\_1 Submission for Distributed Data Analytics Lab - Group 1

This notebook contains explanations for questions in the tutorial and code descriptions as comments. The notebook can be run without any particular configuration. Required explanations and descriptions can be found for each section followed by the code and graphics.

Descriptions and Explanations can be directly accessed by searching for '[EXPLANATION]'.

```
[ ]: # Importing the libraries
import numpy as np
import math
import matplotlib.pyplot as plt
import pandas as pd
```

Matrix Multiplication: Create a numpy matrix A of dimensions  $n \times m$ , where  $n = 100$  and  $m = 20$ . Initialize Matrix A with random values. Create a numpy vector v of dimensions  $m \times 1$ . Initialize the vector v with values from a normal distribution using  $\mu = 2$  and standard deviation = 0.01. Perform the following operations: 2. Iteratively multiply (element-wise) each row of the matrix A with vector v and sum the result of each iteration in another vector c. This operation needs to be done with for-loops, not numpy built-in operations.

```
[ ]: #Setting variables for the data
n = 100
m = 20
u = 2
sigma = 0.01

np.random.seed(0) #setting the seed for the random number generator
A = np.random.randn(n,m) #generating a random matrix
v = np.random.normal(loc=u, scale = sigma, size=(m,1)) #generating a random
↪vector
c = np.zeros((A.shape[0],v.shape[1])) #creating a matrix of zeros
myList=[] #creating an empty list

for i in range(A.shape[0]): #for loop to iterate through the rows of the matrix
    sumBin = 0 # initializing the sum variable to 0
    for j in range(A.shape[1]): #for loop to iterate through the columns of the
↪matrix
        product = A[i,j]*v[j,0] #calculating the product of the row and column
```

```
sumBin = sumBin + product #adding the product to the sum
c[i,0] = sumBin #setting the value of the row to the sum
```


2. Find mean and standard deviation of the new vector c.

```
[ ]: # Mean and standard deviation of the c
print('Mean :', np.mean(c, axis=0), 'Standard Deviation:', np.std(c, axis=0)) #
    ↪printing the mean and standard deviation of c
```

Mean : [-0.63158641] Standard Deviation: [8.64465701]

3. Plot the histogram of vector c using 5 bins.

```
[ ]: # Plotting the histogram
plt.hist(c, bins = 5, color='yellow', edgecolor='black', linewidth=1,
    ↪align='mid') #plotting the histogram
plt.title('Histogram of the vector c')
plt.show()
```



Grading Program: This task puts you in the position that I end up at the end of every semester. Which is, grading your work and issuing the grades. In this task you are required to use the 'Grades.csv' File that has been provided on learnweb.

- Read the data from the csv.
- Compute the sum for all subjects for each student.

```
[ ]: # Grading Program
df = pd.read_csv('Grades.csv') #reading the csv file
collist = df.columns #creating a list of the columns
subList = collist[2:-1] #creating a list of the subjects
calculated_total = df[subList].sum(axis=1) #calculating the total of the
↳subjects
df['Calculated_Totals'] = calculated_total #adding the calculated total to the
↳dataframe
print(df.head()) #printing the first 5 rows of the dataframe
```

	First Name	Last Name	English	Maths	Science	German	Sports	\
0	Robyn	Hobgood	60.95	24.77	20.60	69.32	8.36	
1	Eddy	Swearngin	100.00	12.99	100.00	52.24	100.00	
2	Leoma	Bridgman	83.37	100.00	78.69	100.00	19.50	
3	Arnetta	Pearl	87.75	100.00	86.93	87.90	41.73	
4	Maryland	Colby	100.00	100.00	100.00	18.87	88.72	

	Final Grade	Calculated_Totals
0	184.00	184.00
1	365.23	365.23
2	381.56	381.56
3	404.31	404.31
4	407.59	407.59

- Compute the average of the point for each student. (total points are 500).

```
[ ]: # Calculating Average Point
student_average = calculated_total/500 #calculating the average point
df['Student_Average'] = student_average #adding the average point to the
↳dataframe
print(df.head()) #printing the first 5 rows of the dataframe
```

	First Name	Last Name	English	Maths	Science	German	Sports	\
0	Robyn	Hobgood	60.95	24.77	20.60	69.32	8.36	
1	Eddy	Swearngin	100.00	12.99	100.00	52.24	100.00	
2	Leoma	Bridgman	83.37	100.00	78.69	100.00	19.50	
3	Arnetta	Pearl	87.75	100.00	86.93	87.90	41.73	
4	Maryland	Colby	100.00	100.00	100.00	18.87	88.72	

	Final Grade	Calculated_Totals	Student_Average
0	184.00	184.00	0.36800
1	365.23	365.23	0.73046
2	381.56	381.56	0.76312
3	404.31	404.31	0.80862
4	407.59	407.59	0.81518

- Compute the standard deviation of point for each student.

```
[ ]: # Standard Deviation
# Calculating standard deviation across row for each student
std = df[subList].std(axis=1) #calculating the standard deviation of the
    ↪subjects
df['Standard_Deviation'] = std #adding the standard deviation to the dataframe
print(df.head()) #printing the first 5 rows of the dataframe
```

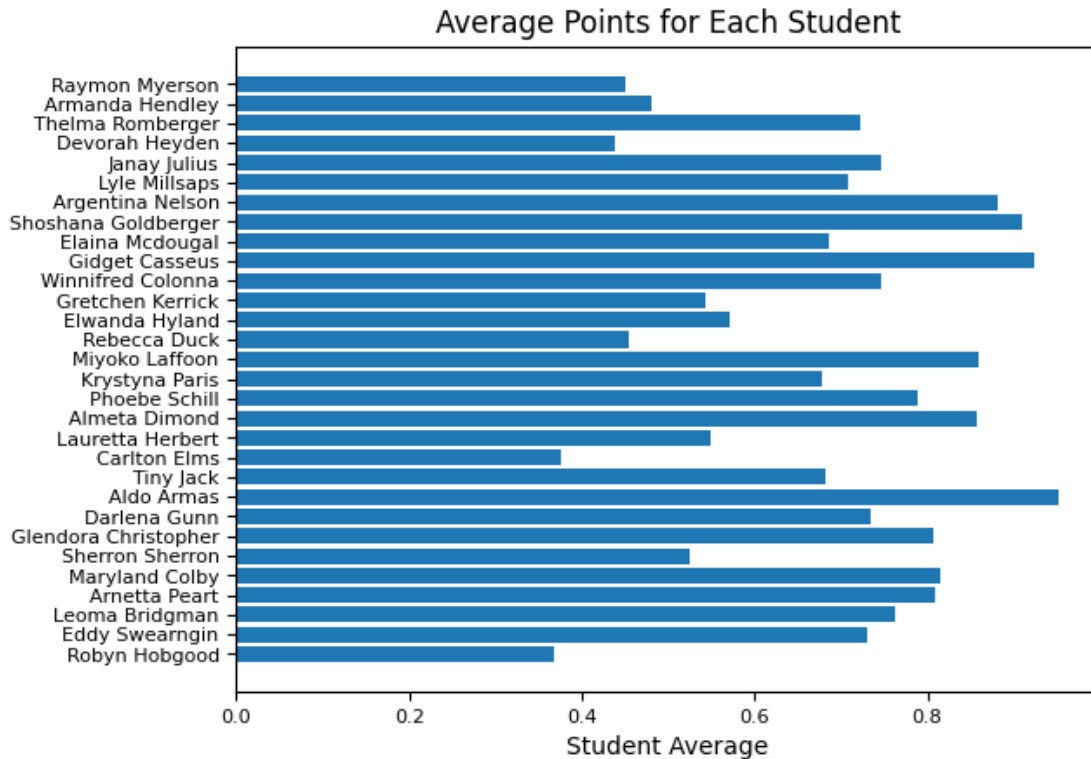
	First Name	Last Name	English	Maths	Science	German	Sports	\
0	Robyn	Hobgood	60.95	24.77	20.60	69.32	8.36	
1	Eddy	Swearngin	100.00	12.99	100.00	52.24	100.00	
2	Leoma	Bridgman	83.37	100.00	78.69	100.00	19.50	
3	Arnetta	Peart	87.75	100.00	86.93	87.90	41.73	
4	Maryland	Colby	100.00	100.00	100.00	18.87	88.72	

	Final Grade	Calculated_Totals	Student_Average	Standard_Deviation
0	184.00	184.00	0.36800	26.724368
1	365.23	365.23	0.73046	39.430848
2	381.56	381.56	0.76312	33.186278
3	404.31	404.31	0.80862	22.535389
4	407.59	407.59	0.81518	35.360266

- Plot the average points for all the students (in one figure).

```
[ ]: # Plot for Average Points
# Using matplotlib code example to plot student average points
# Reference: https://matplotlib.org/stable/gallery/lines_bars_and_markers/
    ↪bar_label_demo.html

nameList=df['First Name'] + ' ' + df['Last Name'] #creating a list of the names
fig, ax = plt.subplots() # creating a figure and axes
ax.barh(nameList, df['Student_Average'], align='center') #plotting the bar chart
ax.set_yticks(nameList, labels=nameList) #setting the y ticks to the names
ax.set_xlabel('Student Average') #setting the x label
ax.set_title('Average Points for Each Student') #setting the title
ax.tick_params(axis='x', labelsize=8) #setting the tick size
ax.tick_params(axis='y', labelsize=8) #setting the tick size
plt.show()
```



- For each student assign a grade based on the following rubric.

```
[ ]: # Grade Assignment
# Grades are assigned to the students by rounding the average to nearest upper_
integer to allow for all grades to be assigned

grades = [] #creating an empty list
gpa = df['Student_Average']*100 #calculating the percentage
for marks in gpa: #for loop to iterate through the student averages
    marks = math.ceil(marks) #rounding the marks to the nearest upper integer
    if marks <= 100 and marks >= 96:
        grades.append('A+')
    elif marks <= 90 and marks >= 90:
        grades.append('A')
    elif marks <= 89 and marks >= 86:
        grades.append('A-')
    elif marks <= 85 and marks >= 80:
        grades.append('B+')
    elif marks <= 79 and marks >= 76:
        grades.append('B')
    elif marks <= 75 and marks >= 70:
        grades.append('B-')
```

```

elif marks <= 69 and marks >= 66:
    grades.append('C+')
elif marks <= 60 and marks >= 65:
    grades.append('C')
elif marks <= 59 and marks >= 56:
    grades.append('D+')
else:
    grades.append('F')

df['Grades'] = grades #adding the grades to the dataframe
print(df.head()) #printing the first 5 rows of the dataframe

```

	First Name	Last Name	English	Maths	Science	German	Sports \
0	Robyn	Hobgood	60.95	24.77	20.60	69.32	8.36
1	Eddy	Swearngin	100.00	12.99	100.00	52.24	100.00
2	Leoma	Bridgman	83.37	100.00	78.69	100.00	19.50
3	Arnetta	Peart	87.75	100.00	86.93	87.90	41.73
4	Maryland	Colby	100.00	100.00	100.00	18.87	88.72


	Final Grade	Calculated_Totals	Student_Average	Standard_Deviation	Grades
0	184.00	184.00	0.36800	26.724368	F
1	365.23	365.23	0.73046	39.430848	B-
2	381.56	381.56	0.76312	33.186278	B
3	404.31	404.31	0.80862	22.535389	B+
4	407.59	407.59	0.81518	35.360266	B+

- Plot the histogram of the final grades

```

[ ]: # Grades Histogram
plt.title('Histogram Grades') #setting the title of the histogram
plt.hist(grades, color='red', edgecolor='black', linewidth=1, align='mid')
    ↪ #plotting the histogram
plt.show() #showing the histogram

```



- Exercise 2: Linear Regression through exact form. (10 Points) In this exercise, you will implement linear regression that was introduced in the introduction Machine Learning Lecture. The method we are implementing here today is for a very basic univariate linear regression.
- Generate 3 sets of simple data, each consisting of a matrix  $A$  with dimensions  $100 \times 2$ . Initialize the sets of data with normal distribution  $\mu = 2$  and  $\sigma = [0.01, 0.1, 1]$  so that each dataset has a different  $\sigma$ . You may assume that the first column of  $A$  represents the predictor data ( $X$ ), whereas the second column of matrix  $A$  represents the target data ( $Y$ ).
- Implement LEARN-SIMPLE-LINREG algorithm and train it using matrix  $A$  to learn values of  $b_0$  and  $b_1$ .
- Implement LEARN-SIMPLE-LINREG algorithm and train it using matrix  $A$  to learn values of  $b_0$  and  $b_1$ .
- The following section define the necessary variable and the functions for this exercise.

```
[ ]: sigmaList = [0.01, 0.1, 1] #creating a list of the standard deviations
rows = 100 #setting the number of rows
cols = 2 #setting the number of columns
u = 2 #setting the mean

def GENERATE_DATA(rows, cols, u, sigma): #function to generate the data
```

```

    np.random.seed(0) #setting the seed for the random number generator for
    ↪reproducibility
    A = np.random.normal(loc=u, scale = sigma, size=(rows, cols)) #generating a
    ↪random matrix
    return A[:,0], A[:,1] #returning the first column and the second column of
    ↪the matrix

def LEARN_SIMPLE_LINREG(X, Y): #function to learn the simple linear regression
    xmean = np.mean(X) #calculating the mean of the x values
    ymean = np.mean(Y) #calculating the mean of the y values
    b1 = sum((X-xmean)*(Y-ymean))/sum((X-xmean)**2) #calculating the slope
    b0 = ymean - (b1 * xmean) #calculating the y-intercept
    return(b0,b1) #returning the y-intercept and the slope

def PREDICT_SIMPLE_LINREG(b0, b1, x): #function to predict the simple linear
    ↪regression
    y = b0 + b1*x #calculating the y value
    return y #returning the y value

```


- [EXPLANATION] Data can be plotted to see how normal distribution looks with increasing the standard deviation values. As evident in the section below, with the increase in the standard deviation value, the spread of the data is increasing further away about the mean point.


```

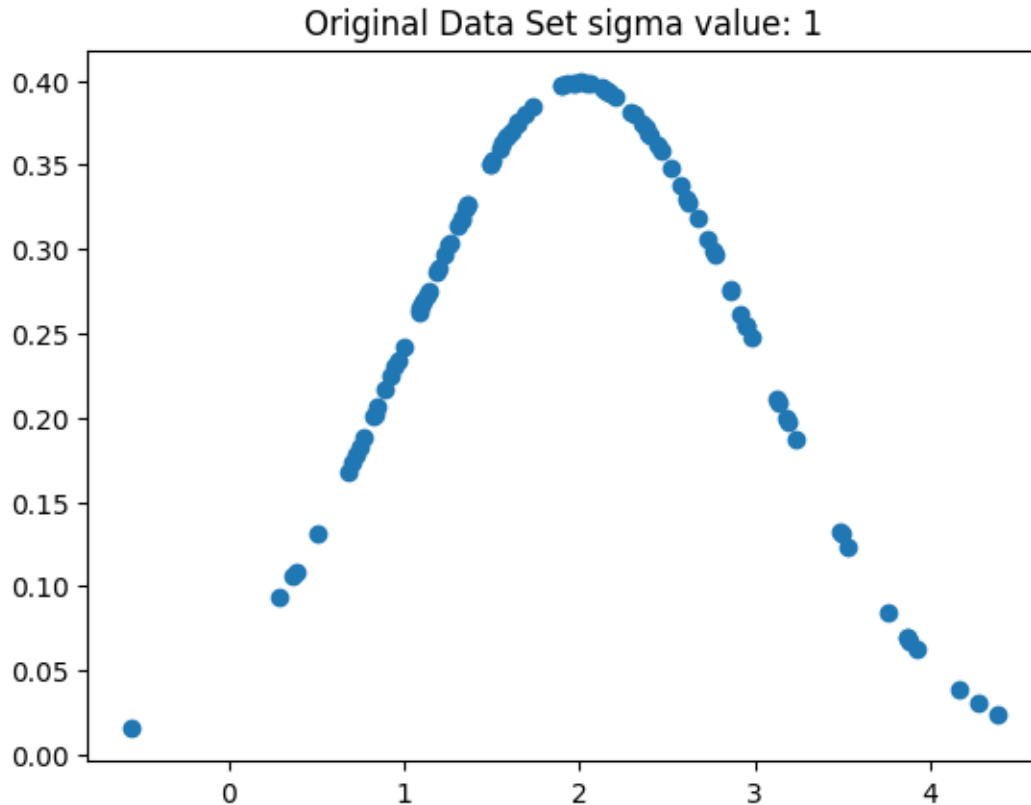
[ ]: # Lets have a look at the data
import scipy.stats as stats
for sig in sigmaList: #for loop to iterate through the standard deviations
    X,Y = GENERATE_DATA(rows, cols, u, sig) #generating the data
    plt.title('Original Data Set'+ ' '+'sigma value: '+str(sig))
    plt.plot(X, stats.norm.pdf(X, u, sig), 'o') #plotting the original data
    plt.show()

```










- [EXPLANATION] In the section below, the earlier implemented algorithms are being called, with a for loop to iterate over the provided std. deviation values. The original data and the predicted data is then presented in the graphs generated for each data set as per the std. deviation value.


```
[ ]: v1=np.zeros((3,2)) #creating a matrix of zeros
      i1 = 0
      for sig in sigmaList: #for loop to iterate through the standard deviations
          X,Y = GENERATE_DATA(rows, cols, u, sig) #generating the data
          b0, b1 = LEARN_SIMPLE_LINREG(X,Y) #learning the simple linear regression
          v1[i1,:] = np.array((b0, b1)) # creating a vector of the y-intercept and
          ↪the slope to compare later with numpy.linalg.lstsq
          print('Sigma:', sig, 'B0:', b0, 'B1:', b1) #printing the slope and
          ↪y-intercept
          y = PREDICT_SIMPLE_LINREG(b0, b1, X) #predicting the y value
          plt.figure() #creating a figure
          plt.title('Data with Line of Estimation'+ ' '+'sigma value: '+str(sig))
          ↪#setting the title
          plt.scatter(X,Y, label='Original Data') #plotting the data
          plt.plot(X,y,'g', label='Best fit line', linewidth=2.4) #plotting the line
          ↪of estimation
```

```
plt.legend() #plotting the legend
plt.show() #showing the figure
i1 += 1
```


Sigma: 0.01 B0: 2.0491493939501475 B1: -0.023860917890473884



Sigma: 0.1 B0: 2.0619974174729507 B1: -0.023860917890475914



Sigma: 1 B0: 2.1904776527009386 B1: -0.023860917890475678




- Put  $b_0$  to zero and rerun the program to generate the predicted line. Comment on the change you see for the varying values of  $\sigma$ .
- [EXPLANATION] As evident in the next section, the value of  $b_0 = 0$ , hence this translates to the y-intercept to be 0. As per the definition of y-intercept, which is the point where the graph intersects the y-axis. Now as evident from the values and the plots, as standard deviation increases the spread of the data is also increasing around the mean point. The relation standard deviation to  $b_1$  is by  $(b_1 = \text{sum}((X - x_{\text{mean}}) * (Y - y_{\text{mean}})) / \text{sum}((X - x_{\text{mean}})^2))$ . Moreover, due to  $b_0$  being 0, that is the intercept, the predicted line is further from the actual data cloud exactly by the shift factor equal to  $b_0$ .


```
[ ]: b0 = 0 # setting the y-intercept to 0
for sig in sigmaList: #for loop to iterate through the standard deviations
    X,Y = GENERATE_DATA(rows, cols, u, sig) #generating the data
    _, b1 = LEARN_SIMPLE_LINREG(X,Y) #learning the simple linear regression
    print('Sigma:', sig, 'B0:', b0, 'B1:', b1) #printing the slope and
    ↪y-intercept
    y = PREDICT_SIMPLE_LINREG(b0, b1, X) #predicting the y value
    plt.figure() #creating a figure
    plt.title('Data with Line of Estimation'+ ' '+'sigma value: '+str(sig))
    ↪#setting the title
    plt.scatter(X,Y, label='Original Data') #plotting the data
```

```
plt.plot(X,y,'g', label='Best fit line') #plotting the best fit line  
plt.legend() #plotting the legend  
plt.show() #showing the plot
```

Sigma: 0.01 B0: 0 B1: -0.023860917890473884




Sigma: 0.1 B0: 0 B1: -0.023860917890475914



Sigma: 1 B0: 0 B1: -0.023860917890475678






- Put  $b_1$  to zero and rerun the program to generate the predicted line. Comment on the change you see for the varying values of  $\sigma$ .
- [EXPLANATION] The equation in our case looks like,  $Y = b_1X + b_0$ . This means if  $b_1 = 0$ , the relationship of  $X$  and  $Y$  is vanished. This means the output from our implementation has no direct correlation other than the value of  $b_0$  to the input, i.e. ( $b_0 = y_{\text{mean}} - (b_1 \cdot x_{\text{mean}})$ ). In such a case varying values of standard deviation have no impact on the value of prediction 'y' since y is not equals to a constant that is  $b_0$ , the equation looks like  $y = b_0$ . Hence, in all three cases there is a line with 0 gradient (straight, horizontal line) that is only determined through a scalar value  $b_0$ .


```
[ ]: b1 = 0 # setting the slope to 0
for sig in sigmaList: #for loop to iterate through the standard deviations
    X,Y = GENERATE_DATA(rows, cols, u, sig) #generating the data
    b0, _ = LEARN_SIMPLE_LINREG(X,Y) #learning the simple linear regression
    print('Sigma:', sig, 'B0:', b0, 'B1:', b1) #printing the slope and
    ↪y-intercept
    y = PREDICT_SIMPLE_LINREG(b0, b1, X) #predicting the y value
    plt.figure() #creating a figure
    plt.title('Data with Line of Estimation'+ ' '+'sigma value: '+str(sig))
    ↪#setting the title
    plt.scatter(X,Y, label='Original Data') #plotting the data
```

```
plt.plot(X,y,'g', label='Best fit line') #plotting the best fit line
plt.legend() #plotting the legend
plt.show() #showing the plot
```


Sigma: 0.01 B0: 2.0491493939501475 B1: 0



Sigma: 0.1 B0: 2.0619974174729507 B1: 0



Sigma: 1 B0: 2.1904776527009386 B1: 0



- Use `numpy.linalg.lstsq` to replace step 2 for learning values of `b0` and `b1` . Explain the difference between your values and the values from `numpy.linalg.lstsq`.
- [EXPLANATION] In this section, first `numpy.linalg.lstsq()` is used then the results obtained from manual implementation and `numpy.linalg.lstsq()` are compared.
- [EXPLANATION] There seems to be negligible difference in the output, which is due to higher decimal positions only, since both of the algorithms use sum of squares to find the error values for the calculation of the gradient and the intercept.

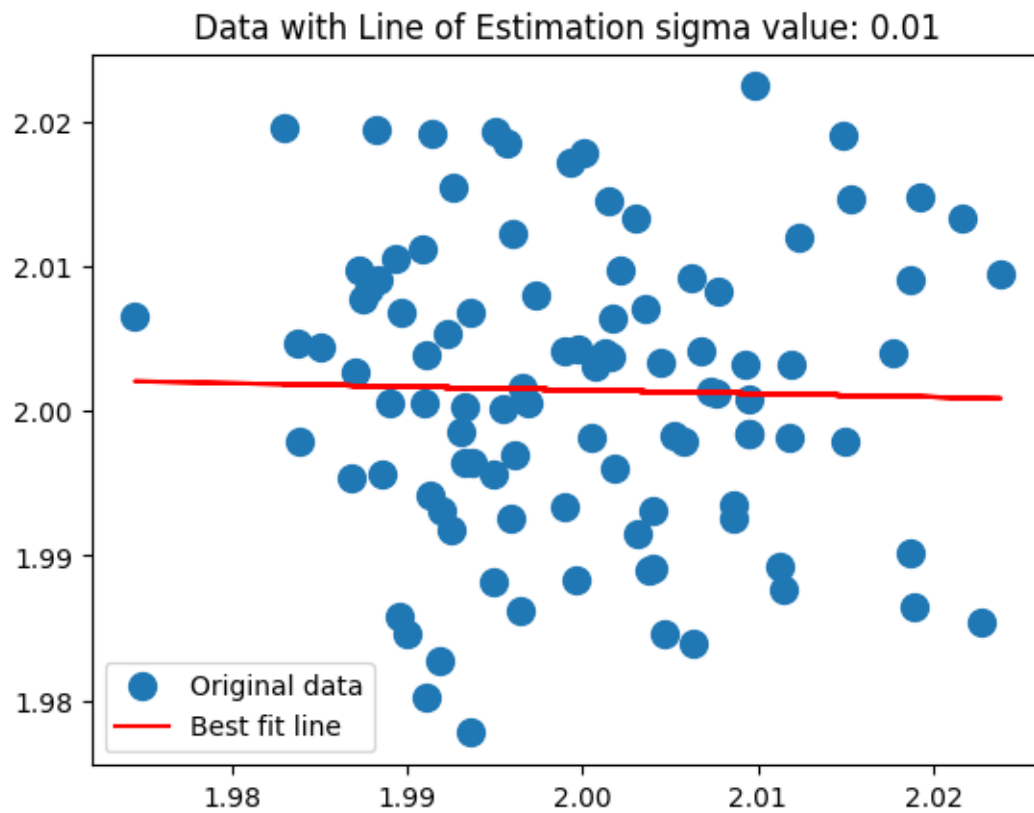
```
[ ]: # In this section we will be using numpy.linalg.lstsq() to solve the linear
      ↪ regression problem.
# As per the documentation for numpy.linalg.lstsq(). It processes the data in
      ↪ the form y=A_dash * p
# where A_dash = [[X, 1]] and p = [m, c] where m is the slope(b1_dash) and
      ↪ c(b0_dash) is the y-intercept.
# In this case we will loop over the standard deviation list, generate A from
      ↪ X, then apply the np.linalg.lstsq()
# Reference: https://numpy.org/doc/stable/reference/generated/numpy.linalg.
      ↪ lstsq.html
v2=np.zeros((3,2)) #creating a matrix of zeros
i2=0 #setting the index to 0
```

```


for sig in sigmaList: #for loop to iterate through the standard deviations
    x,y = GENERATE_DATA(rows, cols, u, sig) #generating the data
    A = np.vstack([x, np.ones(len(x))]).T # using vstack to stack the x values
    ↪and the ones values
    b1_dash, b0_dash = np.linalg.lstsq(A, y, rcond=None)[0] #calculating the
    ↪slope and y-intercept
    v2[i2,:] = np.array((b0_dash, b1_dash)) # creating a vector of the
    ↪y-intercept and the slope to compare with values from custom implementation
    print('Sigma:', sig, 'B0:', b0_dash, 'B1:', b1_dash) #printing the slope
    ↪and y-intercept
    plt.figure() #creating a figure
    plt.title('Data with Line of Estimation'+ ' '+'sigma value: '+str(sig))
    plt.plot(x, y, 'o', label='Original data', markersize=10)
    plt.plot(x, b1_dash*x + b0_dash, 'r', label='Best fit line')
    plt.legend()
    plt.show()
    i2 += 1
v3 = v2-v1 # subtracting the values from the custom implementation from the
    ↪numpy.linalg.lstsq() implementation
v3 = sum((v3)**2) #calculating the sum of the squared values
print('The difference in values from manual implementation and numpy.linalg.
    ↪lstsq() is: ',v3)

```


Sigma: 0.01 B0: 2.049149393950148 B1: -0.023860917890473943



Sigma: 0.1 B0: 2.06199741747295 B1: -0.023860917890474752



Sigma: 1 B0: 2.190477652700939 B1: -0.023860917890475508



The difference in values from manual implementation and `numpy.linalg.lstsq()` is:  
[1.18329136e-30 1.38323900e-30]